

Submission 01 Moritz Perschke

Exercise 1

Task 1

Study how to submit jobs in SLURM, how to check their state and how to cancel them.

This was mentioned in the proseminar: - to submit a job, use **squeue** - to check on a jobs status, the alias **squ** can be used on the cluster, to show the jobs of \$USER - to cancel, use **scancel**

Task 2

Prepare a submission script that starts an arbitrary executable, e.g. `/bin/hostname`

```
#!/bin/bash

# Execute job in the partition "lua" unless you have special requirements.
#SBATCH --partition=lua
# Name your job to be able to identify it later
#SBATCH --job-name test
# Redirect output stream to this file
#SBATCH --output=output.log
# Maximum number of tasks (=processes) to start in total
#SBATCH --ntasks=1
# Maximum number of tasks (=processes) to start per node
#SBATCH --ntasks-per-node=1
# Enforce exclusive node allocation, do not share with other jobs
#SBATCH --exclusive

/bin/hostname
```

The example script does this already.

The output of the script (in output.log in the current dir) is as follows:

```
n001.intern.lcc3.intra.uibk.ac.at
```

Task 3

In your opinion, what are the 5 most important parameters available when submitting a job and why? What are possible settings of these parameters, and what effect do they have?

I think the example script already includes some of the most important/useful parameters. For example exclusive node allocation, number of tasks and the job name, which I see myself using for almost all jobs in this PS.

I think the `--mem-per-cpu` or `--mem-per-gpu` parameter is one of the most important parameters, which is not part of the example script. It might be the case that I know the minimum memory requirement of my job beforehand, and I know that some nodes do not have enough, while others do. In this case this parameter is basically required in order to ensure that my job does not fail.

For a job that reads data from a file, such as measurements, the `--input` parameter is also essential. This parameter connects the scripts `stdin` to this file.

On a larger supercomputer, used by multiple projects at the same time, the `--clusters` parameter might be very important. This parameter specifies the cluster to be used by the job, which could come in very handy in a larger supercomputer. Using this parameter different projects can be assigned different clusters for example, which allows for better sharing of a supercomputer that is meant to be used by different projects at the same time.

For tasks that have specific requirements, the `--constraint` or `-C` parameter can ensure the job only runs on nodes with a specific label. These labels are assigned by the administrator. Examples mentioned in the documentation include `intel` or `amd` (for vendor specific cpus), `graphics` (for a node with a gpu) or a specific rack. So if the supercomputer has nodes with and without gpus, and the job requires a gpu, this parameter can ensure the usage of nodes that can adequately handle the job.

There are also a few parameters that deal with ‘accounting’. For example `--account`, which specifies the account to be charged, might be *very* important on an actual scientific cluster. For the lcc3, which by now only exists for teaching purposes (as far as I understand it), I don’t see the importance. Noone pays on the lcc3 for resources used, so why would I want to change the account that gets charged.

Task 4

How do you run your program in parallel? What environment setup is required?

This was also discussed in the proseminar.

```
#!/bin/bash

# Execute job in the partition "lua" unless you have special requirements.
#SBATCH --partition=lua
# Name your job to be able to identify it later
#SBATCH --job-name test
# Redirect output stream to this file
#SBATCH --output=output.log
# Maximum number of tasks (=processes) to start in total
#SBATCH --ntasks=X
# Maximum number of tasks (=processes) to start per node
```

